

Feature Engineering

1 feature engineering

#label encoding using for different techniques

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
[23]: df=pd.read_csv("train.csv")
```

```
[3]: df.head()
```

```
[3]:      Sex
0    male
1  female
2  female
3  female
4    male
```

```
[4]: # 1.pd.get_dummies
df1=pd.get_dummies(df["Sex"])
```

```
[5]: df1
```

```
[5]:      female  male
0         0     1
1         1     0
2         1     0
3         1     0
4         0     1
..      ...    ...
886       0     1
887       1     0
888       1     0
889       0     1
890       0     1
```

[891 rows x 2 columns]

```
[6]: # 2.Label Encoder  
from sklearn.preprocessing import LabelEncoder
```

```
[9]: oe=LabelEncoder()
```

```
[17]: Label=oe.fit_transform(df["Sex"])
```

```
[19]: Label
```

```
[19]: array([[1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1,  
0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0,  
0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,  
0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1,  
0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0,  
1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1,  
0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1,  
1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1,  
1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1,  
0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0,  
1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1,  
0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1,  
1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1,  
0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1,  
0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0,  
1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1,  
0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1,  
1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0,  
1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1,  
0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 1,  
1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1,  
1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0,  
1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0,  
1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0,  
1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1,  
0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1,  
1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1,  
1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1,  
0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0,  
1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1,  
1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1,  
0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0,  
1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0,
```

```
1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1,
0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0,
0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1])
```

[39]: *# 3. ordinal encoding*

```
from numpy import asarray
from sklearn.preprocessing import OrdinalEncoder
# define data
data = asarray(['red'], ['green'], ['blue']))
print(data)
# define ordinal encoding
encoder = OrdinalEncoder()
# transform data
result = encoder.fit_transform(data)
print(result)
```

```
['red']
['green']
['blue']
[[2.]
 [1.]
 [0.]]
```

[40]: *# 4. one hot encoding*

```
from numpy import asarray
from sklearn.preprocessing import OneHotEncoder
# define data
data = asarray(['red'], ['green'], ['blue']))
print(data)
# define one hot encoding
encoder = OneHotEncoder(sparse=False)
# transform data
onehot = encoder.fit_transform(data)
print(onehot)
```

```
['red']
['green']
['blue']
[[0. 0. 1.]
 [0. 1. 0.]
 [1. 0. 0.]]
```

[41]: **import datetime**

[42]: today_date=datetime.datetime.today()

```

[43]: today_date
[43]: datetime.datetime(2021, 10, 2, 22, 19, 45, 120558)
[44]: today_date-datetime.timedelta(3)
[44]: datetime.datetime(2021, 9, 29, 22, 19, 45, 120558)
[46]: days=[today_date-datetime.timedelta(x) for x in range(0,15)]
days
[46]: [datetime.datetime(2021, 10, 2, 22, 19, 45, 120558),
datetime.datetime(2021, 10, 1, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 30, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 29, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 28, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 27, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 26, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 25, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 24, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 23, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 22, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 21, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 20, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 19, 22, 19, 45, 120558),
datetime.datetime(2021, 9, 18, 22, 19, 45, 120558)]
[47]: import pandas as pd
data=pd.DataFrame(days)
data.columns=["Day"]
[48]: data.head()
[48]:
          Day
0 2021-10-02 22:19:45.120558
1 2021-10-01 22:19:45.120558
2 2021-09-30 22:19:45.120558
3 2021-09-29 22:19:45.120558
4 2021-09-28 22:19:45.120558
[:]: data["weekday"]=data["Day"].dt.weekday_name
data.head()
[52]: dictionary={"Monday":1,"Tuesday":2,"Wednesday":3,"Thursday":4,"Friday":
, 5,"Saturday":6,"Sunday":7

}
[53]: dictionary
[53]: {"Monday": 1,
      "Tuesday": 2,

```

```
'Wednesday': 3,
'Thursday': 4,
'Friday': 5,
'Saturday': 6,
'Sunday': 7}
```

```
[ ]: data["weekday_ordinal"]=data["weekday"].map(dictionary)
[55]: # 5.count fequency
train_set = pd.read_csv("http://archive.ics.uci.edu/ml/
machine-learning-databases/adult/adult.data" , header = None,index_col=None)
train_set.head()

[ ]: columns=[1,3,5,6,7,8,9,13]
[ ]: train_set=train_set[columns]
[ ]: train_set.
[ ]: columns=["Employment", "Degree", "Status", "Designation", "family_job", "Race", "Sex", "Country"]
train_set.head()
[ ]: for feature in train_set.columns[:]:
[ ]:     print(feature,":",len(train_set[feature].unique()),"labels")
country_map=train_set["Country"].value_counts().to_dict()
[ ]: train_set["Country"]=train_set["Country"].map(country_map)
[ ]: train_set.head(20)
```

2 Missing Values- Feature Engineering

3 # Mean/ Median /Mode imputation

When should we apply? Mean/median imputation has the assumption that the data are missing completely at random(MCAR). We solve this by replacing the NAN with the most frequent occurrence of the variables

```
[ ]: #using the median
age=df.fillna(df.Age.median())
df1["Age"]=df1["Age"].apply(np.int64)

[ ]: #using the mean
df["Age"]=df.fillna(df.Age.mean())

[ ]: #When we deal the catogorical values we use mode
df["SibSp"]=df.fillna(df.SibSp.Mode()[0])

[57]: df.isnull().sum()
[57]: PassengerId    0
Survived         0
Pclass           0
```

```
Name          0
Sex           0
Age          177
SibSp        0
Parch        0
Ticket       0
Fare         0
Cabin       687
Embarked     2
dtype: int64
```

```
[59]: df=pd.read_csv('train.csv',usecols=['Age','Fare','Survived'])
df.head()
```

```
[59]:   Survived   Age   Fare
0         0  22.0  7.2500
1         1  38.0 71.2833
2         1  26.0  7.9250
3         1  35.0 53.1000
4         0  35.0  8.0500
```

```
[60]: ## Lets go and see the percentage of missing values
df.isnull().mean()
```

```
[60]: Survived    0.000000
Age          0.198653
Fare          0.000000
dtype: float64
```

```
[61]: def impute_nan(df,variable,median):
        df[variable+"_median"]=df[variable].fillna(median)
```

```
[62]: median=df.Age.median()
median
```

```
[62]: 28.0
```

```
[63]: impute_nan(df,'Age',median)
df.head()
```

```
[63]:   Survived   Age   Fare  Age_median
0         0  22.0  7.2500         22.0
1         1  38.0 71.2833         38.0
2         1  26.0  7.9250         26.0
3         1  35.0 53.1000         35.0
4         0  35.0  8.0500         35.0
```


```
[64]: print(df['Age'].std())
print(df['Age_median'].std())
```

```
14.526497332334044
13.019696550973194
```

```
[65]: import matplotlib.pyplot as plt
      %matplotlib inline

[66]: fig = plt.figure()
      ax = fig.add_subplot(1 1 1)
      df["Age"].plot(kind="kde", ax=ax)
      df.Age_median.plot(kind="kde", ax=ax, color="red")
      lines, labels = ax.get_legend_handles_labels()
      ax.legend(lines, labels, loc="best")

[66]: <matplotlib.legend.Legend at 0x13057ea5470>
```



4 Feature Scaling

#Normalization

Normalization is a scaling technique in which values are shifted and rescaled so that they end up ranging between 0 and 1. It is also known as Min-Max scaling.

Here's the formula for normalization:

Normalization equation

Here, X_{max} and X_{min} are the maximum and the minimum values of the feature respectively.

When the value of X is the minimum value in the column, the numerator will be 0, and hence X' is 0. On the other hand, when the value of X is the maximum value in the column, the numerator is equal to the denominator and thus the value of X' is 1. If the value of X is between the minimum and the maximum value, then the value of X' is between 0 and 1.

#Standardization

Standardization is another scaling technique where the values are centered around the mean with a unit standard deviation. This means that the mean of the attribute becomes zero and the resultant distribution has a unit standard deviation.

Here's the formula for standardization:

Standardization equation

Feature scaling: Muis the mean of the feature values andFeature scaling: Sigma is the standard deviation of the feature values. Note that in this case, the values are not restricted to a particular range.

Now, the big question in your mind must be when should we use normalization and when should we use standardization Let's find out

```
[67]: from sklearn.preprocessing import StandardScaler,MinMaxScaler,RobustScaler
```

```
[ ]: std=StandardScaler()  
std.fit_transform(df["Variable"].reshape(-1,1).values())
```

```
[ ]: min=MinMaxScaler()  
min.fit_transform(["variable"].reshape(0,1).values())
```

```
[ ]: Rob=RobustScaler()  
Rob.fit_transform(df["variable"].reshape(0,1).values())
```

5 #Outliers Tretment

```
[69]: df.head()
```

```
[69]:
```

	Survived	Age	Fare	Age_median
0	0	22.0	7.2500	22.0
1	1	38.0	71.2833	38.0
2	1	26.0	7.9250	26.0
3	1	35.0	53.1000	35.0
4	0	35.0	8.0500	35.0


```
[70]: df["Age"].isnull().sum()
```

```
[70]: 177
```

```
[71]: import seaborn as sns
```


```
[72]: sns.distplot(df["Age"].dropna())
```

```
[72]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```


```
[73]: sns.distplot(df['Age'].fillna(100))
```

```
[73]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```




```
[74]: figure=df.Age.hist(bins=50)  
      figure.set_title("Age")  
      figure.set_xlabel("Age")  
      figure.set_ylabel("No of passenger")
```

```
[74]: Text(0, 0.5, "No of passenger")
```



```
[77]: figure=df.boxplot(column="Age")
```



```
[76]: df["Age"].describe()
```

```
[76]: count    714.000000
      mean      29.699118
      std       14.526497
      min        0.420000
      25%       20.125000
      50%       28.000000
      75%       38.000000
      max       80.000000
      Name: Age, dtype: float64
```

6 If The Data Is Normally Distributed We use this

```
[78]: ### Assuming Age follows A Gaussian Distribution we will calculate the  
      boundaries which differentiates the outliers
```

```
upper_boundary=df["Age"].mean() + 3* df["Age"].std()  
lower_boundary=df["Age"].mean() - 3* df["Age"].std()  
print(lower_boundary), print(upper_boundary),print(df["Age"].mean())
```


```
-13.88037434994331  
73.27860964406095  
29.69911764705882
```

[78]: (None, None, None)

7 If Features Are Skewed We Use the below Technique


```
[79]: figure=df.Fare.hist(bins=50)
      figure.set_title("Fare")
      figure.set_xlabel("Fare")
      figure.set_ylabel("No of passenger")
```

[79]: Text(0, 0.5, 'No of passenger')



```
[80]: df.boxplot(column="Fare")
```

[80]: <AxesSubplot:>



```
[81]: df["Fare"].describe()
```

```
[81]: count    891.000000
      mean     32.204208
      std      49.693429
      min       0.000000
      25%      7.910400
      50%     14.454200
      75%     31.000000
      max     512.329200
      Name: Fare, dtype: float64
```

```
[82]: ## Lets compute the Interquantile range to calculate the boundaries
      IQR=df.Fare.quantile(0.75)-df.Fare.quantile(0.25)
```

```
[83]: lower_bridge=df["Fare"].quantile(0.25)-(IQR*1.5)
      upper_bridge=df["Fare"].quantile(0.75)+(IQR*1.5)
      print(lower_bridge), print(upper_bridge)
```

```
-26.724
65.6344
```

```
[83]: (None, None)
```

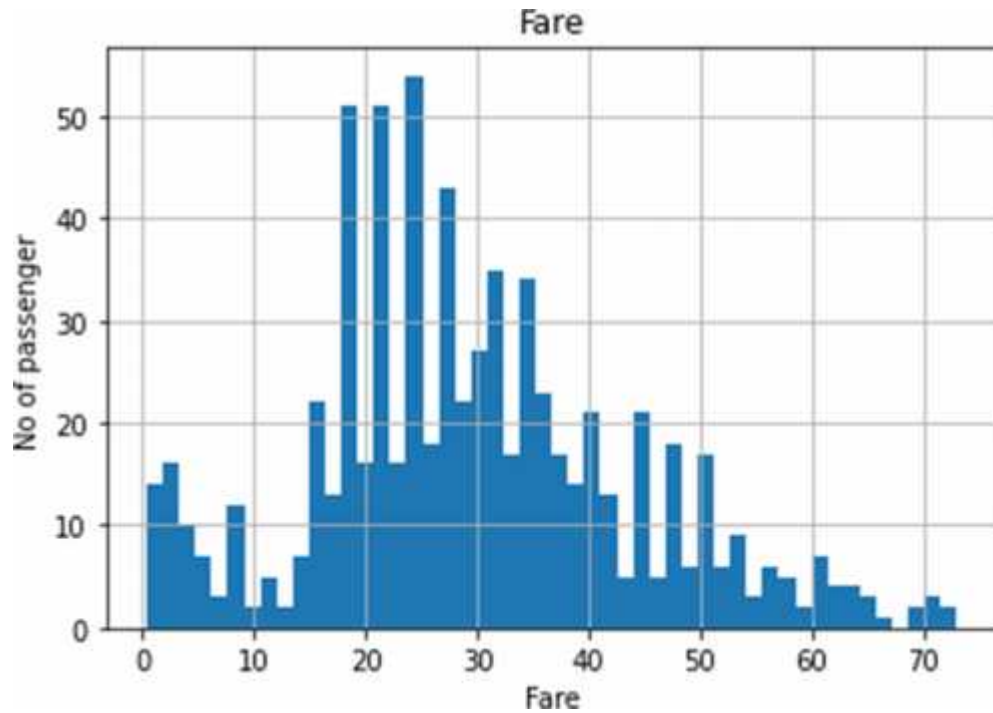
```
[84]: data=df.copy()
```

```
[85]: data.loc[data["Age"]>=73,"Age"]=73
```

```
[86]: data.loc[data["Fare"]>=100,"Fare"]=100
```

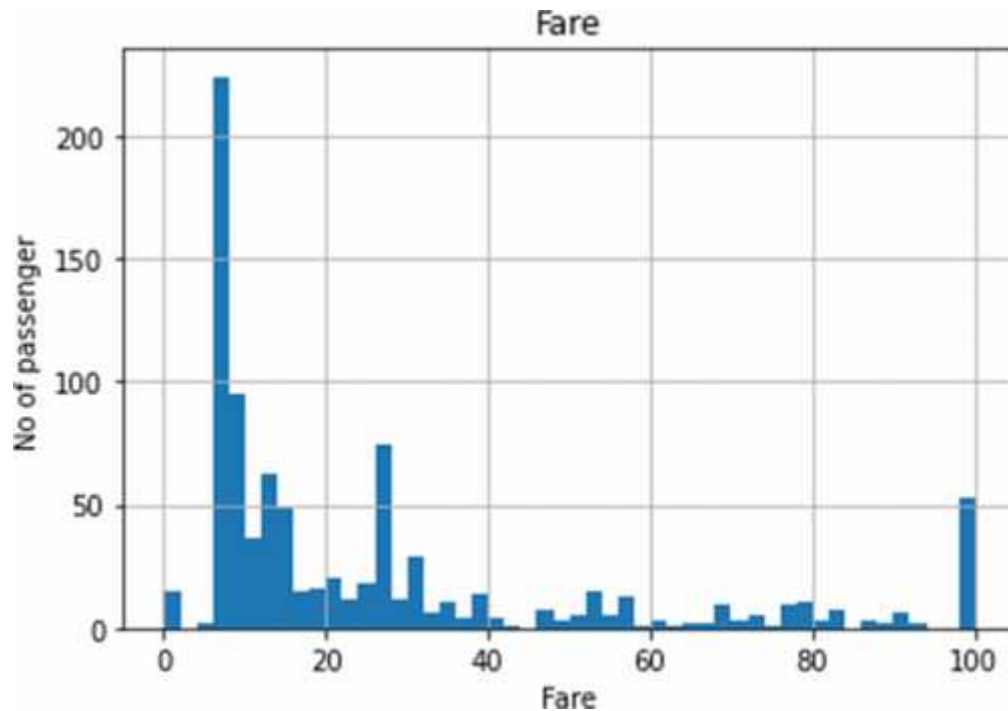
```
[87]: figure=data.Age.hist(bins=50)
figure.set_title("Fare")
figure.set_xlabel("Fare")
figure.set_ylabel("No of passenger")
```

```
[87]: Text(0, 0.5, "No of passenger")
```



```
[88]: figure=data.Fare.hist(bins=50)
figure.set_title("Fare")
figure.set_xlabel("Fare")
figure.set_ylabel("No of passenger")
```

```
[88]: Text(0, 0.5, "No of passenger")
```



8 Feature Selection

```
[ ]: # Remove The correlated
threshold=0.8

[ ]: # 1. find and remove correlated features

def correlation(dataset, threshold):
    col_corr = set() # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in
absolute coeff value
                colname = corr_matrix.columns[i] # getting the name of column
                col_corr.add(colname)
    return col_corr

[ ]: correlation(df.iloc[:, :-1], threshold)

[ ]: # 2.information gain get from independent variable and dependent variable

from sklearn.feature_selection import mutual_info_classif
```

```

[: mutual_info=mutual_info_classif(X,y)
[: mutual_data=pd.Series(mutual_info,index=X.columns)
mutual_data.sort_values(ascending=False)
[: # 3. Tree classifier extract the important features from dataset

from sklearn.ensemble import ExtraTreeClassifier
from sklearn.datasets import make_classification
X, y = make_classification(n_features=4, random_state=0)
clf = ExtraTreesClassifier(n_estimators=100, random_state=0)
clf.fit(X, y)

[: # 4. chi2 feature selection
from sklearn.datasets import load_iris
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2

# Load iris data
iris_dataset = load_iris()

# Create features and target
X = iris_dataset.data
y = iris_dataset.target

# Convert to categorical data by converting data to integers
X = X.astype(int)

# Two features with highest chi-squared statistics are selected
chi2_features = SelectKBest(chi2, k = 2)
X_kbest_features = chi2_features.fit_transform(X, y)

```